

LLM Workflow Release Ops

Phil Stevens

philipstevens.github.io

Operate releases with eval gating, drift detection, and rollback discipline so you ship safely over time.

SCOPE

One or more production LLM workflows requiring ongoing release discipline and drift management.

PRIMARY OUTPUT

Gated releases + drift monitoring + threshold maintenance + rollback readiness (sustained over time).

Best Fit

- Your workflow is live, and regressions or drift are a recurring risk.
- You want repeatable releases with clear evidence and fast rollback.
- You already have (or want to establish) an eval harness and baseline gates.

Inputs Required

INPUT	MINIMUM	WHY IT MATTERS
Baseline eval suite	Runnable gate suite with must-pass subset	Release ops is about preventing regressions over time.
Telemetry access	Version identifiers in trace/logs; production signals	Without this, incident response becomes guesswork.
Release process	Canary/shadow/ramp capability (even manual)	Allows safe rollout and rapid rollback.
Owner plus cadence	Someone who can approve releases and review weekly	Ensures accountability and keeps drift monitoring actionable.

If you don't have these yet, the first month often starts with a short Build & Harden sprint to establish them.

How the Work Runs

CADENCE	WHAT HAPPENS
WEEKLY	Review production signals, sample quality on a rotating set, triage drift indicators and prioritize fixes.
PER RELEASE	Run eval gates on proposed changes, analyze deltas vs baseline, define rollout plan and rollback triggers.
MONTHLY	Update test sets from incidents, review and ratchet gate thresholds, calibrate alerts to reduce noise.

Deliverables

ARTIFACT	PURPOSE
Maintained eval suite (test sets updated from incidents)	Keeps coverage current as product and failure modes evolve
Release gate reports	Documents what changed, what passed/failed, and ship decision
Weekly stability memos	Summarizes trends, drift signals, and recommended fixes
Drift monitoring dashboards and alerts	Surfaces regressions before users report them
Updated thresholds and gate configurations	Ratchets quality bar as workflow stabilizes
Rollback runbook (tested and current)	Ensures fast recovery when releases fail

Definition of Done

- ✓ Every release has a gate report (what changed, what passed/failed, decision recorded).
- ✓ Drift and regression monitoring is in place with clear signals and response steps.
- ✓ Thresholds and test sets are maintained (updates are versioned and reviewed).
- ✓ Rollback is documented and exercised so failures are recoverable quickly.

Boundaries

- Building brand-new workflows from scratch (handled via Audit + Build & Harden).
- Large re-architecture or platform build work.
- Large labeling efforts or full-time QA programs.

Next Step

Review one workflow's current gates and monitoring, then define the first month's reliability plan. If you're pre-prod, start with an Audit or Build & Harden sprint.

Book an intro call:

calendly.com/philipstevens4/intro



Scan to book a call

Email: philipstevens4@gmail.com

Web: philipstevens.github.io