

LLM Workflow Production Readiness Checklist

Version 1.0 | Philip Stevens | philipstevens4@gmail.com

How to Use This Checklist

This checklist is designed for engineering teams preparing to ship an LLM-powered workflow to production. Work through each section before release. Items marked with [CRITICAL] are non-negotiable; shipping without them significantly increases the risk of production incidents.

Part 1: Eval Gates

Before any release, run these evaluations and verify thresholds are met.

1.1 Accuracy & Quality

Check	Your Threshold	Measured	Pass?
[CRITICAL] Task accuracy on golden set ($n \geq 50$)	___ %	___ %	<input type="checkbox"/>
Accuracy on edge cases subset	___ %	___ %	<input type="checkbox"/>
Accuracy on adversarial/malformed inputs	___ %	___ %	<input type="checkbox"/>
Human preference score (if applicable)	___ /5	___ /5	<input type="checkbox"/>

Setting thresholds: Start with your current baseline. If you don't have one, run the eval and use current performance minus a small buffer (e.g., if you measure 94%, set threshold at 92%). Raise the bar over time.

1.2 Safety & Compliance

Check	Threshold	Measured	Pass?
[CRITICAL] Refusal rate on unsafe inputs	100%	___ %	<input type="checkbox"/>
[CRITICAL] No PII leakage on test set	0 instances	___	<input type="checkbox"/>
Hallucination rate (factual claims)	\leq ___ %	___ %	<input type="checkbox"/>
Compliance with domain-specific rules	100%	___ %	<input type="checkbox"/>

Unsafe input test set: Include prompt injections, attempts to extract system prompts, requests for harmful content, and attempts to bypass guardrails. Minimum 20 cases; 50+ recommended.

1.3 Performance & Cost

Check	Threshold	Measured	Pass?
Latency p50	\leq ____s	____s	<input type="checkbox"/>
[CRITICAL] Latency p95	\leq ____s	____s	<input type="checkbox"/>
Latency p99	\leq ____s	____s	<input type="checkbox"/>
Cost per request (avg)	\leq \$ ____	\$ ____	<input type="checkbox"/>
Token efficiency (output/input ratio)	\leq ____	—	<input type="checkbox"/>

Latency measurement: Measure end-to-end, not just model call time. Include retrieval, preprocessing, validation, and any retries.

1.4 Regression Check

Check	Threshold	Measured	Pass?
[CRITICAL] Regression suite pass rate	100%	____%	<input type="checkbox"/>
No new failures on previously-passing cases	0	—	<input type="checkbox"/>
Performance delta vs. previous version	\leq ____%	____%	<input type="checkbox"/>

Part 2: Failure Mode Coverage

Verify you have detection and mitigation for each failure mode category.

2.1 Output Quality Failures

Failure Mode	Detection Method	Mitigation	<input type="checkbox"/>
Hallucinated facts	Citation verification, factual consistency check	Ground with retrieved docs, add confidence thresholds	<input type="checkbox"/>
Incomplete output	Required field validation, length checks	Structured output schema, retry logic	<input type="checkbox"/>
Wrong format	Schema validation, regex checks	Strict output parsing, fallback formatting	<input type="checkbox"/>
Inconsistent with context	Semantic similarity to input, contradiction detection	Re-ranking, chain-of-thought verification	<input type="checkbox"/>
Outdated information	Timestamp checks on retrieved content	Source freshness filters, recency weighting	<input type="checkbox"/>

2.2 Safety Failures

Failure Mode	Detection Method	Mitigation	<input type="checkbox"/>
Prompt injection executed	Input classification, output anomaly detection	Input sanitization, output filtering, system prompt hardening	<input type="checkbox"/>
PII in output	Regex + NER detection on outputs	PII scrubbing layer, training data audit	<input type="checkbox"/>
Harmful content generated	Content classification on outputs	Output filtering, refusal training	<input type="checkbox"/>
System prompt leaked	Pattern matching for prompt fragments	Instruction hierarchy, output filtering	<input type="checkbox"/>
Unauthorized capability use	Action logging, capability boundaries	Explicit allow-lists, confirmation steps	<input type="checkbox"/>

2.3 Reliability Failures

Failure Mode	Detection Method	Mitigation	<input type="checkbox"/>
Model API timeout	Request timing, circuit breaker triggers	Timeouts, retries with backoff, fallback responses	<input type="checkbox"/>
Rate limit exceeded	429 response tracking	Request queuing, rate limiting at app layer	<input type="checkbox"/>
Context window exceeded	Token counting before calls	Truncation strategy, summarization, chunking	<input type="checkbox"/>
Retrieval returned no results	Empty result detection	Fallback to broader query, graceful degradation	<input type="checkbox"/>
Retrieval returned irrelevant results	Relevance scoring threshold	Re-ranking, score cutoffs, "I don't know" responses	<input type="checkbox"/>

2.4 Upstream Dependency Failures

Failure Mode	Detection Method	Mitigation	
Model behavior changed (silent update)	Eval suite drift detection, output distribution monitoring	Version pinning where possible, automated regression alerts	<input type="checkbox"/>
Embedding model changed	Similarity score distribution shift	Re-index on change, version tracking	<input type="checkbox"/>
Vector DB unavailable	Health checks, latency monitoring	Caching layer, graceful degradation	<input type="checkbox"/>
Source data stale or missing	Freshness checks, data pipeline monitoring	Staleness alerts, fallback sources	<input type="checkbox"/>

Part 3: Release Decision Framework

3.1 Ship / No-SHIP Criteria

SHIP if all of the following are true:

- All **[CRITICAL]** eval gates pass
- No regressions on the regression suite
- All failure modes have detection or mitigation in place
- Rollback tested and verified working
- Monitoring and alerting configured
- Required sign-offs collected

NO-SHIP if any of the following are true:

- Any **[CRITICAL]** eval gate fails
- New regression introduced
- Unmitigated high-severity failure mode discovered
- Rollback not tested or broken
- Missing required sign-off

3.2 Release Artifacts Checklist

Before release, verify these artifacts exist and are versioned:

Artifact	Location	Version	Verified?
Prompt(s)	_____	v_____	<input type="checkbox"/>
System configuration	_____	v_____	<input type="checkbox"/>
Model identifier	_____	_____	<input type="checkbox"/>
Eval suite	_____	v_____	<input type="checkbox"/>
Regression test set	_____	v_____	<input type="checkbox"/>
Retrieval index (if applicable)	_____	v_____	<input type="checkbox"/>

3.3 Rollback Verification

Check	Status
Previous version artifacts accessible	<input type="checkbox"/>
Rollback procedure documented	<input type="checkbox"/>
Rollback tested in staging	<input type="checkbox"/>
Rollback time estimate: _____ minutes	<input type="checkbox"/>
Rollback owner identified: _____	<input type="checkbox"/>

Part 4: Post-Deploy Monitoring

4.1 Real-Time Signals

Configure alerts for these signals before going live:

Signal	Alert Threshold	Current Value	Configured?
Error rate (5xx, exceptions)	> ___ %	___ %	<input type="checkbox"/>
Latency p95	> ___ s	___ s	<input type="checkbox"/>
Request volume anomaly	± ___ % from baseline	___	<input type="checkbox"/>
Cost per hour	> \$ ___	\$ ___	<input type="checkbox"/>
Empty/null response rate	> ___ %	___ %	<input type="checkbox"/>

4.2 Quality Monitoring (Sampled)

Signal	Sample Rate	Check Frequency	Configured?
Human review of random outputs	___ %	Daily / Weekly	<input type="checkbox"/>
Automated quality scoring	___ %	Continuous	<input type="checkbox"/>
User feedback/thumbs tracking	100%	Continuous	<input type="checkbox"/>
Hallucination spot-check	___ %	Daily / Weekly	<input type="checkbox"/>

4.3 Drift Detection

Signal	Detection Method	Check Frequency	Configured?
Output length distribution	Statistical test on rolling window	Daily	<input type="checkbox"/>
Output sentiment/tone	Classifier on sampled outputs	Daily	<input type="checkbox"/>
Refusal rate	Threshold on rolling average	Continuous	<input type="checkbox"/>
Latency trend	Regression on 7-day window	Daily	<input type="checkbox"/>
Eval score trend	Weekly eval run, track over time	Weekly	<input type="checkbox"/>

Part 5: Regression Harness Structure

5.1 Test Case Categories

A complete regression suite should include cases from each category:

Category	Description	Minimum Cases	Your Count
Golden set	Representative inputs with verified correct outputs	50	—
Edge cases	Boundary conditions, unusual but valid inputs	20	—
Adversarial	Prompt injections, malformed inputs, attack attempts	20	—
Historical failures	Cases that broke in previous versions	All	—
High-stakes	Cases where errors have significant consequences	10	—

5.2 Test Case Structure

Each test case should include:

```
{
  "id": "unique-identifier",
  "category": "golden|edge|adversarial|regression|high-stakes",
  "input": { ... },
  "expected_output": { ... } | null,
  "evaluation": {
    "method": "exact_match|semantic_similarity|llm_judge|custom",
    "threshold": 0.95,
    "custom_evaluator": "path/to/evaluator" | null
  },
  "metadata": {
    "added_date": "2024-01-15",
    "source": "production_failure|synthetic|user_reported",
    "severity": "critical|high|medium|low",
    "notes": "..."
  }
}
```

5.3 Harness Requirements

Requirement	Implementation	Done?
Single command to run full suite	<code>make eval</code> or equivalent	<input type="checkbox"/>
Parallelized execution	Configurable concurrency	<input type="checkbox"/>
Deterministic where possible	Fixed seeds, temperature=0	<input type="checkbox"/>
Results persisted	Database or versioned files	<input type="checkbox"/>
Diff against previous run	Automated comparison	<input type="checkbox"/>
CI/CD integration	Runs on PR, blocks on failure	<input type="checkbox"/>
Human-readable report	Summary + drill-down	<input type="checkbox"/>

Part 6: Quick Reference

Red Flags That Should Block Release

1. **Regression on any previously-passing test case** — Something broke
2. **Safety eval failure** — Non-negotiable
3. **Latency p95 above threshold** — Will affect users
4. **Untested rollback** — You will need it eventually
5. **"We'll fix it after launch"** — You probably won't

Common Mistakes

Mistake	Why It Hurts	What to Do Instead
Testing only happy paths	Real traffic includes edge cases and adversarial inputs	Build adversarial test set from day one
Threshold set to current performance	Any variance causes false failures	Set threshold below current with small buffer
Eval suite in notebook, not CI	Gets skipped under deadline pressure	Integrate into PR workflow from start
No rollback testing	Rollback fails when you need it most	Test rollback monthly, after every infra change
Ignoring cost until bill arrives	Budget surprises, rushed optimization	Track cost per request from day one
"Model X is better" without eval	Vibes don't catch regressions	Always run full eval before switching

First 24 Hours Post-Deploy

Hour	Action
0-1	Watch error rate, latency, request volume
1-4	Spot-check 10 random outputs manually
4-8	Review any user feedback/complaints
8-24	Compare quality metrics to pre-deploy baseline
24+	Run full eval suite, compare to release eval

Getting Help

If you're preparing an LLM workflow for production and want expert help with defining acceptance criteria, building eval suites, hardening workflows, or setting up release gates and monitoring:

Book an intro call: calendly.com/philipstevens4/intro

This checklist is provided as a starting point. Adapt thresholds, categories, and checks to your specific workflow and domain requirements.